

```

from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Dense, Dropout, Activation
from keras.utils import np_utils

```

```

(X_train, y_train), (X_test, y_test) = mnist.load_data()

```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist-11493376/11490434 [=====] - 0s 0us/step
11501568/11490434 [=====] - 0s 0us/step

```



```

# let's print the shape before we reshape and normalize
print("X_train shape", X_train.shape)
print("y_train shape", y_train.shape)
print("X_test shape", X_test.shape)
print("y_test shape", y_test.shape)

```

```

# building the input vector from the 28x28 pixels
X_train = X_train.reshape(60000, 784)
X_test = X_test.reshape(10000, 784)
X_train = X_train.astype('float32')
X_test = X_test.astype('float32')

```

```

# normalizing the data to help with the training
X_train /= 255
X_test /= 255

```

```

# print the final input shape ready for training
print("Train matrix shape", X_train.shape)
print("Test matrix shape", X_test.shape)

```

```

X_train shape (60000, 28, 28)
y_train shape (60000,)
X_test shape (10000, 28, 28)
y_test shape (10000,)
Train matrix shape (60000, 784)
Test matrix shape (10000, 784)

```

```

n_classes = 10
print("Shape before one-hot encoding: ", y_train.shape)
Y_train = np_utils.to_categorical(y_train, n_classes)
Y_test = np_utils.to_categorical(y_test, n_classes)
print("Shape after one-hot encoding: ", Y_train.shape)

```

```

Shape before one-hot encoding: (60000,)
Shape after one-hot encoding: (60000, 10)

```

```

model = Sequential()
model.add(Dense(512, input_shape=(784,)))
model.add(Activation('relu'))

```

```
model.add(Dropout(0.2))

model.add(Dense(512))
model.add(Activation('relu'))
model.add(Dropout(0.2))

model.add(Dense(10))
model.add(Activation('softmax'))

model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

history = model.fit(X_train, Y_train,
                    batch_size=128, epochs=20,
                    verbose=2,
                    validation_data=(X_test, Y_test))
```

Epoch 1/20  
469/469 - 10s - loss: 0.2477 - accuracy: 0.9255 - val\_loss: 0.1104 - val\_accuracy: 0.9255  
Epoch 2/20  
469/469 - 13s - loss: 0.1012 - accuracy: 0.9695 - val\_loss: 0.0840 - val\_accuracy: 0.9695  
Epoch 3/20  
469/469 - 8s - loss: 0.0727 - accuracy: 0.9770 - val\_loss: 0.0741 - val\_accuracy: 0.9770  
Epoch 4/20  
469/469 - 7s - loss: 0.0555 - accuracy: 0.9826 - val\_loss: 0.0822 - val\_accuracy: 0.9826  
Epoch 5/20  
469/469 - 7s - loss: 0.0466 - accuracy: 0.9850 - val\_loss: 0.0789 - val\_accuracy: 0.9850  
Epoch 6/20  
469/469 - 8s - loss: 0.0396 - accuracy: 0.9868 - val\_loss: 0.0686 - val\_accuracy: 0.9868  
Epoch 7/20  
469/469 - 7s - loss: 0.0349 - accuracy: 0.9890 - val\_loss: 0.0711 - val\_accuracy: 0.9890  
Epoch 8/20  
469/469 - 7s - loss: 0.0281 - accuracy: 0.9904 - val\_loss: 0.0656 - val\_accuracy: 0.9904  
Epoch 9/20  
469/469 - 7s - loss: 0.0262 - accuracy: 0.9910 - val\_loss: 0.0731 - val\_accuracy: 0.9910  
Epoch 10/20  
469/469 - 7s - loss: 0.0264 - accuracy: 0.9912 - val\_loss: 0.0654 - val\_accuracy: 0.9912  
Epoch 11/20  
469/469 - 8s - loss: 0.0228 - accuracy: 0.9922 - val\_loss: 0.0748 - val\_accuracy: 0.9922  
Epoch 12/20  
469/469 - 8s - loss: 0.0229 - accuracy: 0.9921 - val\_loss: 0.0732 - val\_accuracy: 0.9921  
Epoch 13/20  
469/469 - 8s - loss: 0.0192 - accuracy: 0.9933 - val\_loss: 0.0845 - val\_accuracy: 0.9933  
Epoch 14/20  
469/469 - 8s - loss: 0.0192 - accuracy: 0.9933 - val\_loss: 0.0755 - val\_accuracy: 0.9933  
Epoch 15/20  
469/469 - 7s - loss: 0.0203 - accuracy: 0.9937 - val\_loss: 0.0671 - val\_accuracy: 0.9937  
Epoch 16/20  
469/469 - 7s - loss: 0.0154 - accuracy: 0.9948 - val\_loss: 0.0713 - val\_accuracy: 0.9948  
Epoch 17/20  
469/469 - 8s - loss: 0.0155 - accuracy: 0.9948 - val\_loss: 0.0735 - val\_accuracy: 0.9948  
Epoch 18/20  
469/469 - 7s - loss: 0.0172 - accuracy: 0.9942 - val\_loss: 0.0821 - val\_accuracy: 0.9942  
Epoch 19/20  
469/469 - 7s - loss: 0.0150 - accuracy: 0.9949 - val\_loss: 0.0843 - val\_accuracy: 0.9949  
Epoch 20/20  
469/469 - 8s - loss: 0.0174 - accuracy: 0.9944 - val\_loss: 0.0720 - val\_accuracy: 0.9944

history.history

```
{'loss': [0.24767357110977173,
0.10116829723119736,
0.07266171276569366,
0.055463239550590515,
0.046619877219200134,
0.03961757570505142,
0.034902382642030716,
0.028110342100262642,
0.02616359107196331,
0.026391195133328438,
0.022804247215390205,
0.022861214354634285,
0.01921638660132885,
0.019159788265824318,
0.020306752994656563,
0.015421883203089237,
0.015452152118086815,
0.01722538098692894,
0.014976977370679379,
0.017397306859493256],
'accuracy': [0.9254833459854126,
0.9695000052452087,
0.9769999980926514,
0.9825999736785889,
0.9850333333015442,
0.9868166446685791,
0.9889500141143799,
0.9904166460037231,
0.991016685962677,
0.9912333488464355,
0.9922166466712952,
0.9921333193778992,
0.9933000206947327,
0.9932666420936584,
0.9937333464622498,
0.9947999715805054,
0.9948333501815796,
0.9942166805267334,
0.9949166774749756,
0.9943666458129883],
'val_loss': [0.11038917303085327,
0.08398347347974777,
0.07411109656095505,
0.08219324052333832,
0.07891156524419785,
0.06864918768405914,
0.07106873393058777,
0.06564640998840332,
0.07306645810604095,
0.06542835384607315,
0.07484312355518341,
0.0731651708483696,
0.08449767529964447,
0.0754641443490982,
0.06707073748111725,
0.07127709686756134,
0.07352504879236221,
```

0.08205705881118774,

```
import matplotlib.pyplot as plt
fig = plt.figure()
plt.subplot(2,1,1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='lower right')

plt.subplot(2,1,2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper right')

plt.tight_layout()
```

